

SHEN ZHENG

• (+1)412-909-6416 • shenzhen@andrew.cmu.edu • Personal Webpage • GitHub • Leetcode • Google Scholar

EDUCATION

Carnegie Mellon University

PhD in Robotics | Advisor: [Dr. Srinivasa Narasimhan](#)

Aug. 2024 — Dec. 2027 (Expected)

Pittsburgh, PA

Carnegie Mellon University

Master of Science in Computer Vision (MSCV) | GPA: 4.07/4.33 | Advisor: [Dr. Srinivasa Narasimhan](#)

- Coursework (A+): Learning-based Image Synthesis, Visual Learning and Recognition, Computer Vision, MultiModal Machine Learning, etc.

Aug. 2022 — Dec. 2023

Pittsburgh, PA

Wenzhou-Kean University

Bachelor of Arts in Mathematical Sciences | Major GPA: 3.94/4.00 | Advisor: [Dr. Gaurav Gupta](#)

- Honors: Dean's List (Top 3%); Zhejiang Provincial Government Scholarship (Top 3%); Outstanding Graduate (Top 10%).

Sep. 2017 – Jun. 2021

Wenzhou, China

PROFESSIONAL EXPERIENCES

Waymo

- Perception Software Engineer *Intern* focused on online vectorized HD map construction with long- and short-term memory fusion (*in progress*).
- Developed image-map fusion (+7.5% mAP), ego-pose correction (+1.0% mAP), and global instance ID integration (+1.1% mAP) on Argoverse 2.

May. 2025 – Aug. 2025

Lucid Motors

- Perception Software Engineer in the ADAS perception team responsible for auto-parking, traffic light detection, and blockage detection.
- Improved BEVFormer for auto-parking (reverse & parallel) by using extrinsic calibration to interpolate and smooth edges to enhance curb detection.
- Trained YOLO6 on traffic light images and finetuned arrow types, confidence, IoU, and area thresholds, improving 42+% mAP for night images.
- Developed a binary segmentation model based on CenterNet to detect blockages (ice, snow, mud, blur, raindrops, sun glares), achieving 93%+ IoU.

Feb. 2024 – Jul. 2024

SELECTED PUBLICATIONS

- Anurag Ghosh, **Shen Zheng**, et al., *ROADWork Dataset: Learning to Recognize, Observe, Analyze and Drive Through Work Zones*. **ICCV 2025**.
- Shen Zheng**, et al., *Instance-Warp: Saliency Guided Image Warping for Unsupervised Domain Adaptation*. **WACV 2025**.
- Shen Zheng**, et al., *TPSeNCE: Towards Artifact-Free Realistic Rain Generation for Deraining and Object Detection in Rain*. **WACV 2024**.
- Shen Zheng**, et al., *PointNorm: Dual Normalization is All You Need for Point Cloud Analysis*. **IJCNN 2023** (Oral).
- Shen Zheng** and Gaurav Gupta, *Semantic-Guided Zero-Shot Learning for Low-Light Image/Video Enhancement*. **WACV 2022**.
- Changjie Lu, **Shen Zheng**, et al., *AS-IntroVAE: Adversarial Similarity Distance Makes Robust IntroVAE*. **ACML 2022**.

RESEARCH EXPERIENCES

ROADWork Dataset: Learning to Recognize, Observe, Analyze and Drive Through Work Zones [\[Webpage\]](#)

Feb. 2024 – Nov. 2024

- Proposed ROADWork, the largest open-source work zone dataset, to enhance baseline and foundation models in understanding work zones.
- Achieved +12.8x work zone discovery, +26.2% mAP for object detection, +36.7 SPICE for image captioning, and -21.3% error for path prediction.

Instance-Warp: Saliency Guided Image Warping for Unsupervised Domain Adaptation [\[Webpage\]](#)

May. 2023 – Jan. 2024

- Propose an instance-level image warping to oversample objects, enhancing object feature learning amidst dominant and highly varied backgrounds.
- Integrate image warping and feature unwarping into unsupervised domain adaptation in a task-agnostic way without test-time warping.
- Achieved +6.1% mAP50 for BDD100K Clear → DENSE Foggy object detection and +6.3% mIoU for Cityscapes → ACDC semantic segmentation.

TPSeNCE: Towards Artifact-Free Realistic Rain Generation for Deraining and Object Detection in Rain [\[Webpage\]](#)

Dec. 2022 – Apr. 2023

- Introduce a Triangular Probability Similarity (TPS) loss based on GAN discriminator to minimize the artifacts and distortions during rain generation.
- Propose a Semantic Noise Contrastive Estimation (SeNCE) strategy based on patch similarity and semantic similarity to optimize the rain amount.
- Obtained the best results for rain generation (KID=72.19), rain removal (MUSIQ=61.853), and object detection (mAP50=52.6%) on BDD100K.

PointNorm: Dual Normalization is All You Need for Point Cloud Analysis [\[Webpage\]](#)

Jun. 2022 – Aug. 2022

- Introduced PointNorm, a point cloud analysis network leveraging local mean and global standard deviation to boost loss and gradient stability.
- Proposed a plug-and-play DualNorm that normalize sampled and grouped points to each other to optimize the point cloud density.
- Achieved superior results on shape classification at ScanObjectNN (OA = 86.8%), part segmentation at ShapeNetPart (Inst. mIoU = 86.2%).

Semantic-Guided Zero-Shot Learning for Low-Light Image/Video Enhancement [\[Webpage\]](#)

Apr. 2021 – Jun. 2021

- Proposed a semantic-guided, zero-shot, low-light image enhancement network that consolidates high-level semantics into low-level enhancement.
- Constructed a recurrent image enhancement network that only demands an enhancement factor map with five non-reference loss functions.
- Attained the state-of-the-art UNIQUE/BRISQUE (0.805/27.01), mIOU/mPA (65.87%/74.50%) and inference time (0.001s).

ACADEMIC ACTIVITIES

- Reviewer:** AAAI 2022, IJCNN (2023-2025), ECCV 2024, WACV (2023-2025), CVPR 2025, ICCV 2025, TNNLS, IJCV.
- Co-Instructor:** MATH 3291/3292: *Computer Vision*, Wenzhou-Kean University, Summer 2023 [\[Slides\]](#) [\[Recordings\]](#).
- Guest Speaker:** *Image Processing with Machine Learning*, Fudan University, Spring 2021 [\[Slides\]](#).
- LeetCode Content Creator:** Created 90+ LeetCode tutorial videos on YouTube, garnering over 40,000 views [\[Link\]](#).

SKILLS

- Languages:** Python, C++, Java, R, Matlab, Shell, HTML, LaTeX, Markdown
- Frameworks & Tools:** PyTorch, TensorFlow, Keras, MMCV, Docker, Git, ONNX, CUDA, Jupyter, Ubuntu
- Libraries:** Scikit-Learn, NumPy, OpenCV, SciPy, Pandas, Matplotlib, Seaborn
- DL Models:** CNN, RNN, GAN, VAE, Diffusion, Transformer